

Forouzan

Chapter 23

Process-to-Process Delivery: UDP, TCP, and SCTP

23-1 PROCESS-TO-PROCESS DELIVERY

The transport layer is responsible for process-toprocess delivery—the delivery of a packet, part of a message, from one process to another. Two processes communicate in a client/server relationship, as we will see later.

Topics discussed in this section:

Client/Server Paradigm
Multiplexing and Demultiplexing
Connectionless Versus Connection-Oriented Service
Reliable Versus Unreliable
Three Protocols



The transport layer is responsible for process-to-process delivery.

Figure 23.1 Types of data deliveries

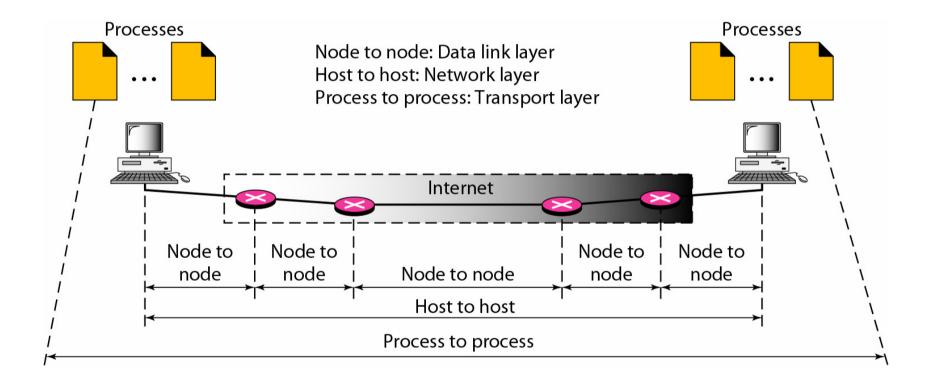


Figure 23.2 Port numbers

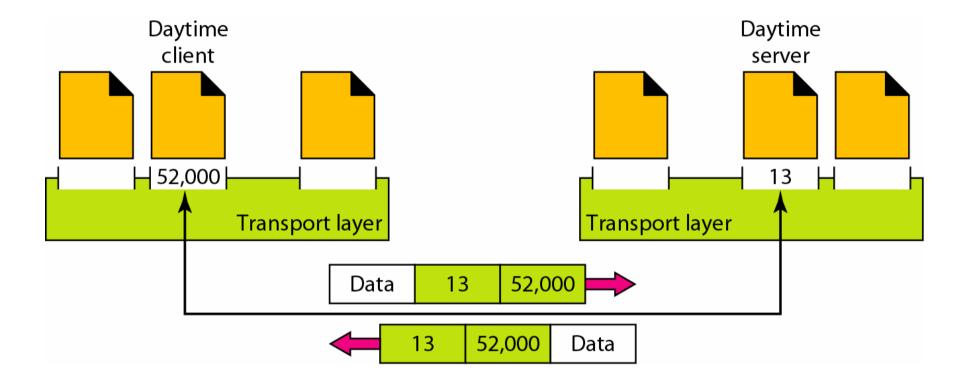


Figure 23.3 IP addresses versus port numbers

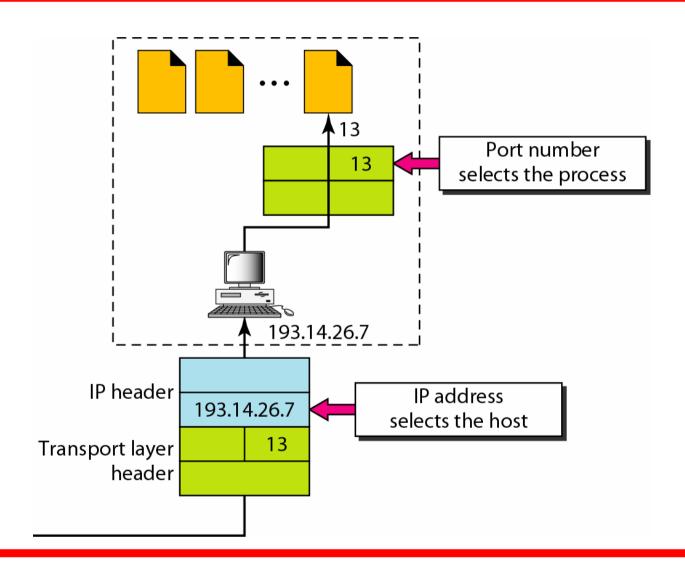


Figure 23.4 IANA ranges

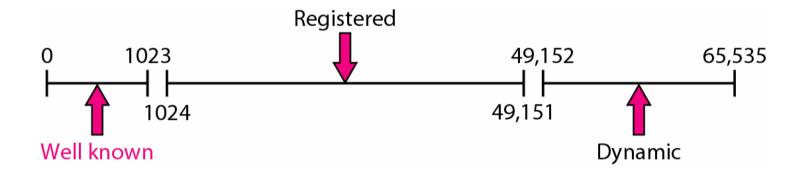


Figure 23.5 Socket address

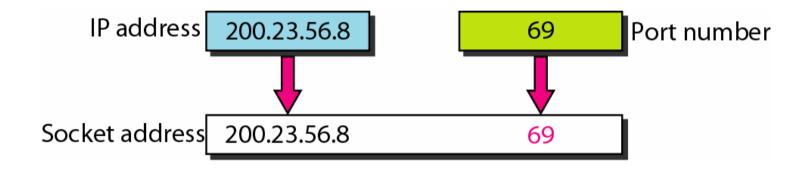


Figure 23.6 Multiplexing and demultiplexing

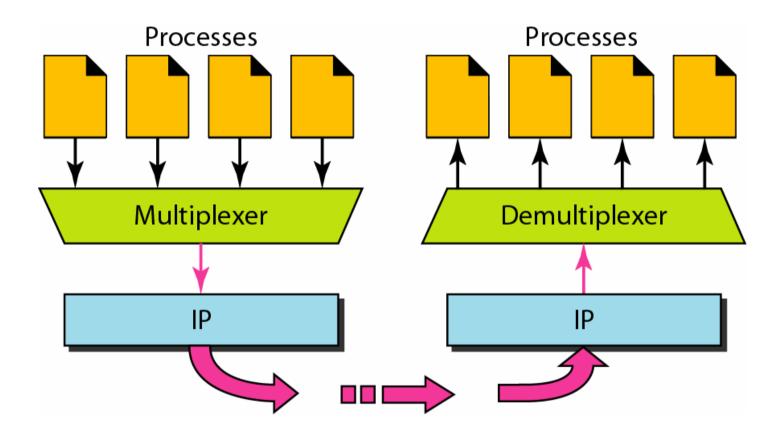


Figure 23.7 Error control

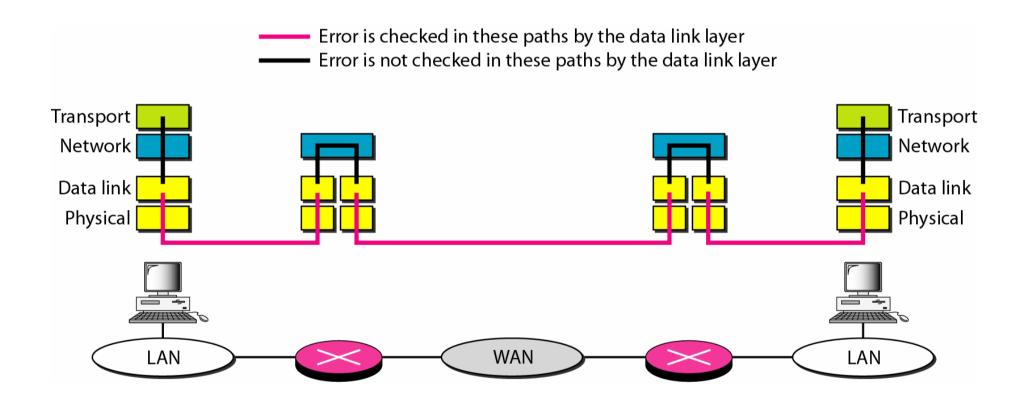
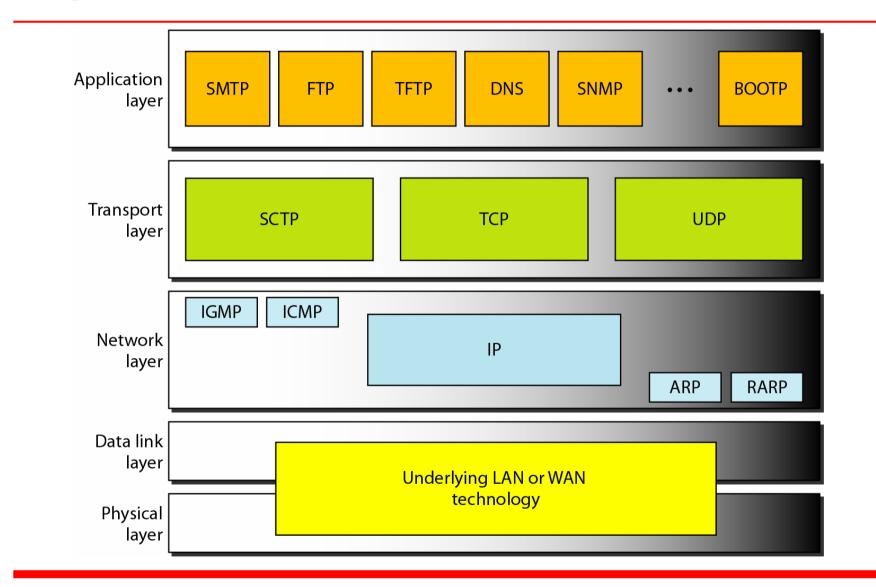


Figure 23.8 Position of UDP, TCP, and SCTP in TCP/IP suite



23-2 USER DATAGRAM PROTOCOL (UDP)

The User Datagram Protocol (UDP) is called a connectionless, unreliable transport protocol. It does not add anything to the services of IP except to provide process-to-process communication instead of host-to-host communication.

Topics discussed in this section:

Well-Known Ports for UDP
User Datagram
Checksum
UDP Operation
Use of UDP

Table 23.1 Well-known ports used with UDP

Port	Protocol	Description
7	Echo	Echoes a received datagram back to the sender
9	Discard	Discards any datagram that is received
11	Users	Active users
13	Daytime	Returns the date and the time
17	Quote	Returns a quote of the day
19	Chargen	Returns a string of characters
53	Nameserver	Domain Name Service
67	BOOTPs	Server port to download bootstrap information
68	ВООТРс	Client port to download bootstrap information
69	TFTP	Trivial File Transfer Protocol
111	RPC	Remote Procedure Call
123	NTP	Network Time Protocol
161	SNMP	Simple Network Management Protocol
162	SNMP	Simple Network Management Protocol (trap)



Example 23.1

In UNIX, the well-known ports are stored in a file called /etc/services. Each line in this file gives the name of the server and the well-known port number. We can use the grep utility to extract the line corresponding to the desired application. The following shows the port for FTP. Note that FTP can use port 21 with either UDP or TCP.

\$ grep	ftp	/etc/services
ftp	21/	tcp
ftp	21/	'udp

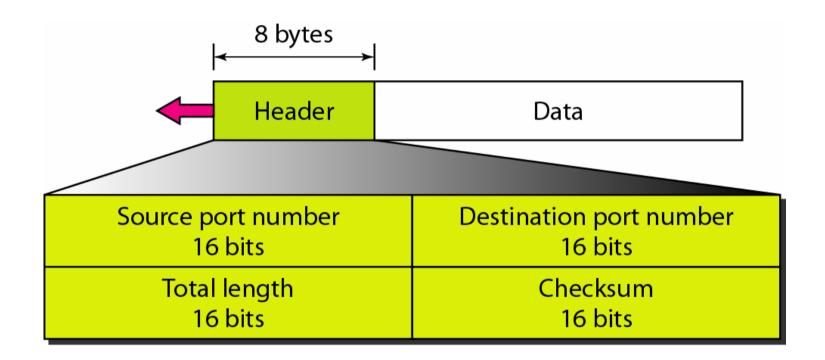


Example 23.1 (continued)

SNMP uses two port numbers (161 and 162), each for a different purpose, as we will see in Chapter 28.

\$ grep	snmp /etc/services	
snmp	161/tcp	#Simple Net Mgmt Proto
snmp	161/udp	#Simple Net Mgmt Proto
snmptrap	162/udp	#Traps for SNMP

Figure 23.9 User datagram format

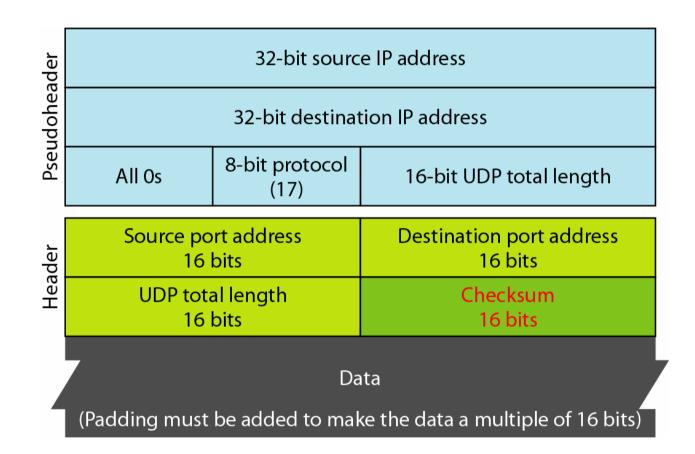




UDP length

= IP length – IP header's length

Figure 23.10 Pseudoheader for checksum calculation



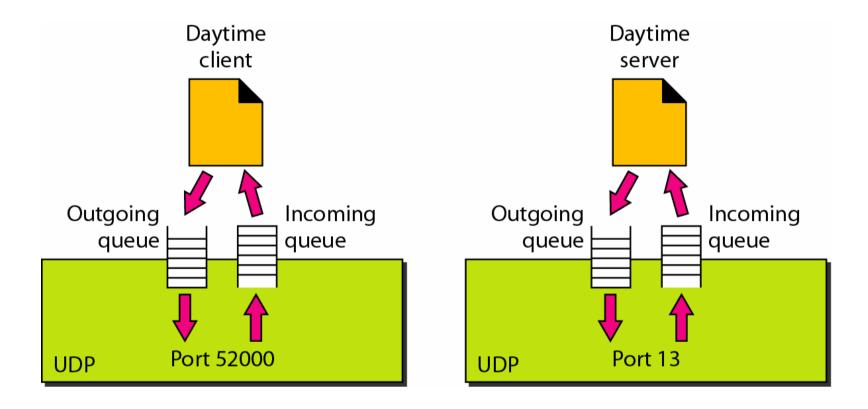
Example 23.2

Figure 23.11 shows the checksum calculation for a very small user datagram with only 7 bytes of data. Because the number of bytes of data is odd, padding is added for checksum calculation. The pseudoheader as well as the padding will be dropped when the user datagram is delivered to IP.

Figure 23.11 Checksum calculation of a simple UDP user datagram

153.18.8.105			
171.2.14.10			
All Os	17	15	
1087		13	
15		All Os	
Т	E	S	Т
I	N	G	All Os

Figure 23.12 Queues in UDP



23-3 TCP

TCP is a connection-oriented protocol; it creates a virtual connection between two TCPs to send data. In addition, TCP uses flow and error control mechanisms at the transport level.

Topics discussed in this section:

TCP Services
TCP Features
Segment
A TCP Connection
Flow Control

Error Control

Table 23.2 Well-known ports used by TCP

Port	Protocol	Description
7	Echo	Echoes a received datagram back to the sender
9	Discard	Discards any datagram that is received
11	Users	Active users
13	Daytime	Returns the date and the time
17	Quote	Returns a quote of the day
19	Chargen	Returns a string of characters
20	FTP, Data	File Transfer Protocol (data connection)
21	FTP, Control	File Transfer Protocol (control connection)
23	TELNET	Terminal Network
25	SMTP	Simple Mail Transfer Protocol
53	DNS	Domain Name Server
67	ВООТР	Bootstrap Protocol
79	Finger	Finger
80	HTTP	Hypertext Transfer Protocol
111	RPC	Remote Procedure Call

Figure 23.13 Stream delivery

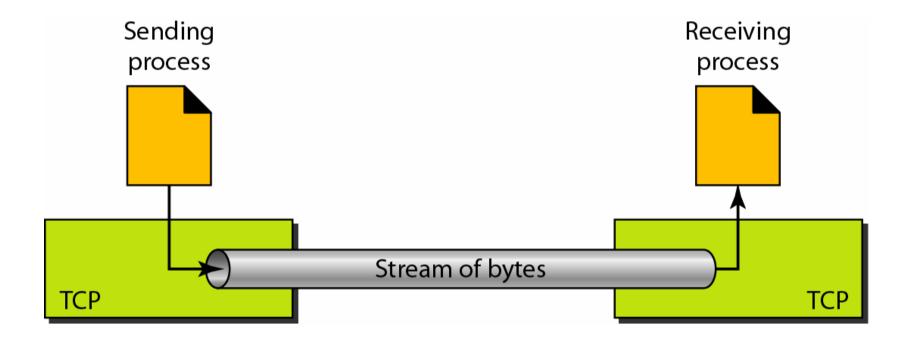


Figure 23.14 Sending and receiving buffers

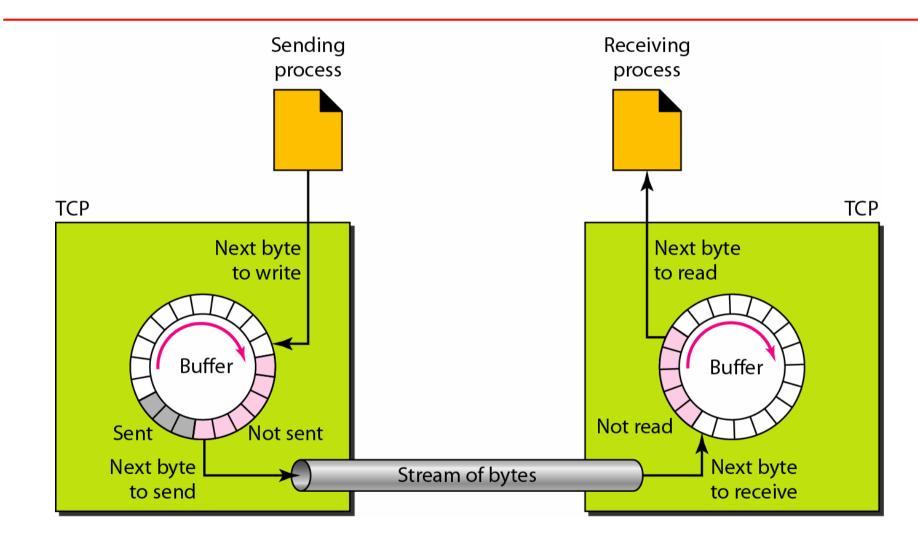
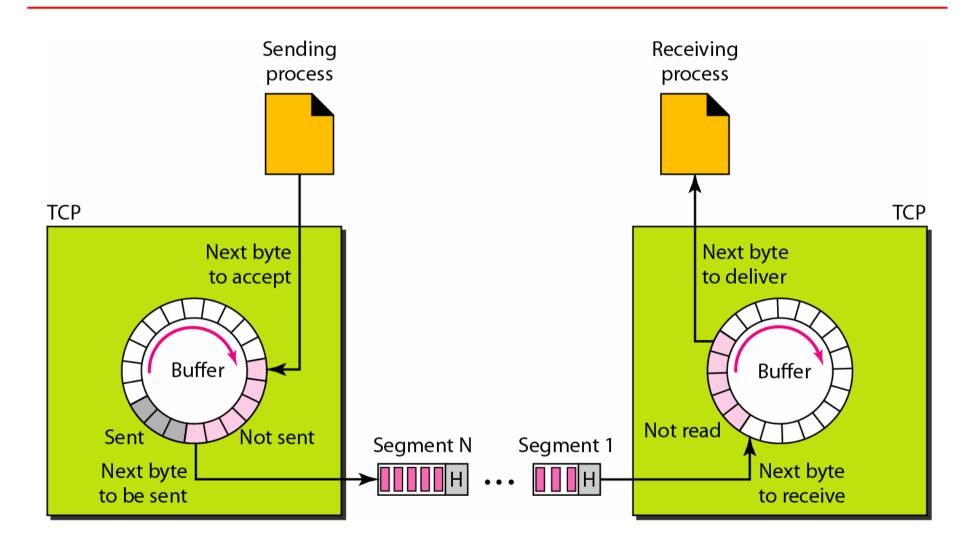


Figure 23.15 TCP segments





The bytes of data being transferred in each connection are numbered by TCP. The numbering starts with a randomly generated number.

Example 23.3

The following shows the sequence number for each segment:

```
      Segment 1
      →
      Sequence Number: 10,001 (range: 10,001 to 11,000)

      Segment 2
      →
      Sequence Number: 11,001 (range: 11,001 to 12,000)

      Segment 3
      →
      Sequence Number: 12,001 (range: 12,001 to 13,000)

      Segment 4
      →
      Sequence Number: 13,001 (range: 13,001 to 14,000)

      Segment 5
      →
      Sequence Number: 14,001 (range: 14,001 to 15,000)
```



The value in the sequence number field of a segment defines the number of the first data byte contained in that segment.



The value of the acknowledgment field in a segment defines the number of the next byte a party expects to receive.

The acknowledgment number is cumulative.

Figure 23.16 TCP segment format

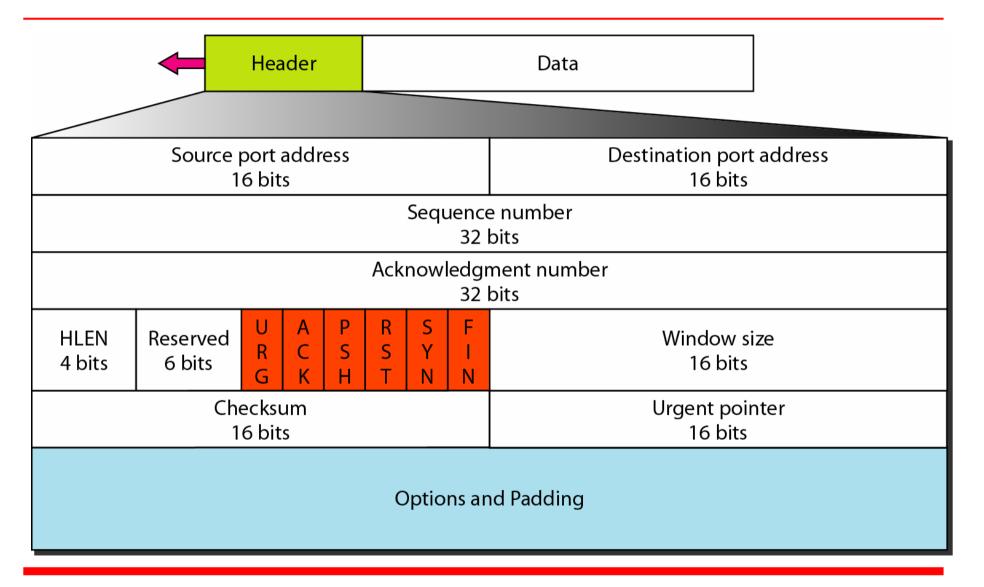


Figure 23.17 Control field

URG: Urgent pointer is valid

ACK: Acknowledgment is valid

PSH: Request for push

RST: Reset the connection

SYN: Synchronize sequence numbers

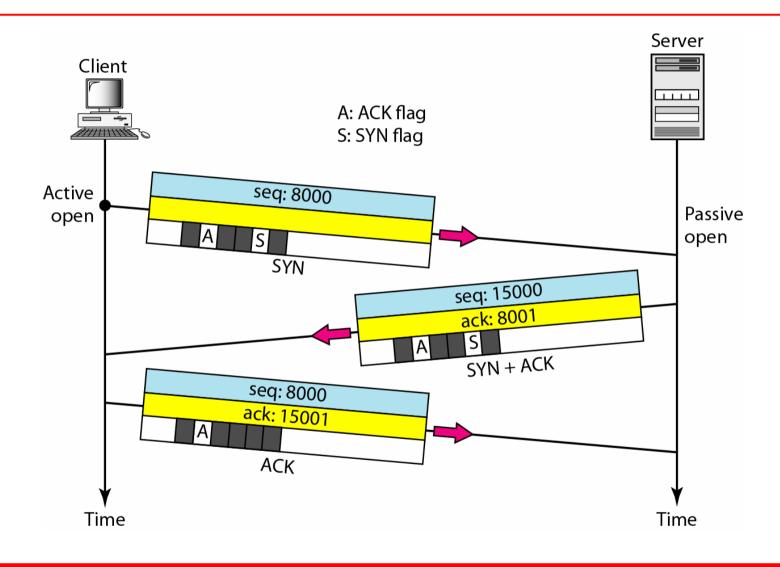
FIN: Terminate the connection

URG ACK PSH RST SYN FIN

Table 23.3 Description of flags in the control field

Flag	Description
URG	The value of the urgent pointer field is valid.
ACK	The value of the acknowledgment field is valid.
PSH	Push the data.
RST	Reset the connection.
SYN	Synchronize sequence numbers during connection.
FIN	Terminate the connection.

Figure 23.18 Connection establishment using three-way handshaking





A SYN segment cannot carry data, but it consumes one sequence number.



A SYN + ACK segment cannot carry data, but does consume one sequence number.



An ACK segment, if carrying no data, consumes no sequence number.

Figure 23.19 Data transfer

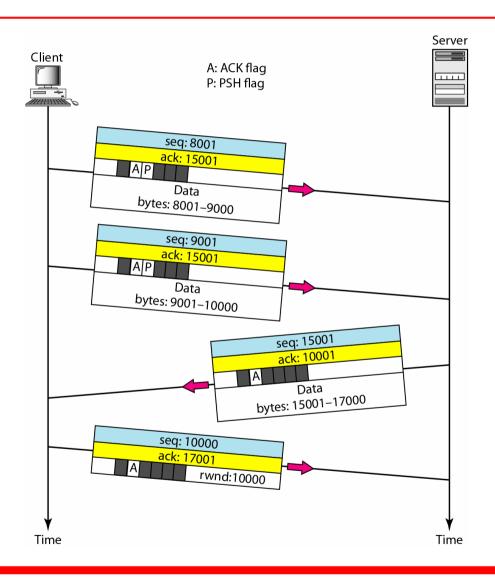
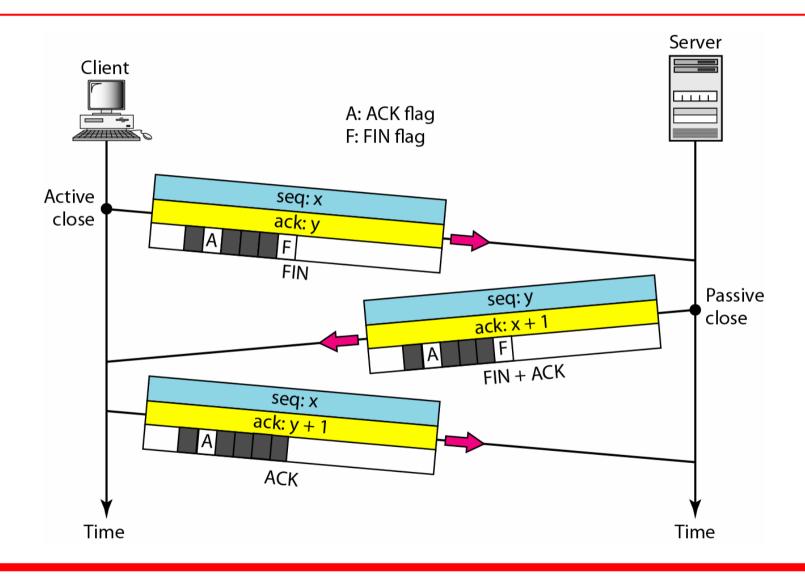


Figure 23.20 Connection termination using three-way handshaking





The FIN segment consumes one sequence number if it does not carry data.



The FIN + ACK segment consumes one sequence number if it does not carry data.

Figure 23.21 Half-close

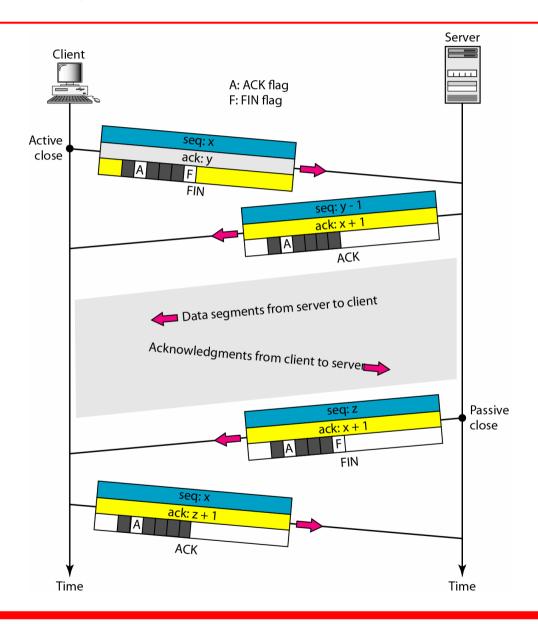
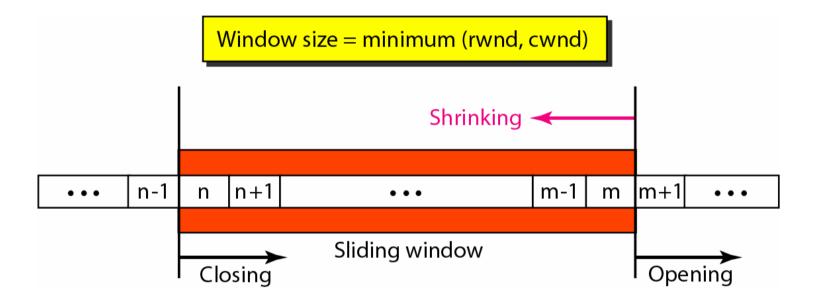


Figure 23.22 Sliding window





A sliding window is used to make transmission more efficient as well as to control the flow of data so that the destination does not become overwhelmed with data.

TCP sliding windows are byte-oriented.

Example 23.4

What is the value of the receiver window (rwnd) for host A if the receiver, host B, has a buffer size of 5000 bytes and 1000 bytes of received and unprocessed data?

Solution

The value of rwnd = 5000 - 1000 = 4000. Host B can receive only 4000 bytes of data before overflowing its buffer. Host B advertises this value in its next segment to A.

Example 23.5

What is the size of the window for host A if the value of rwnd is 3000 bytes and the value of cwnd is 3500 bytes?

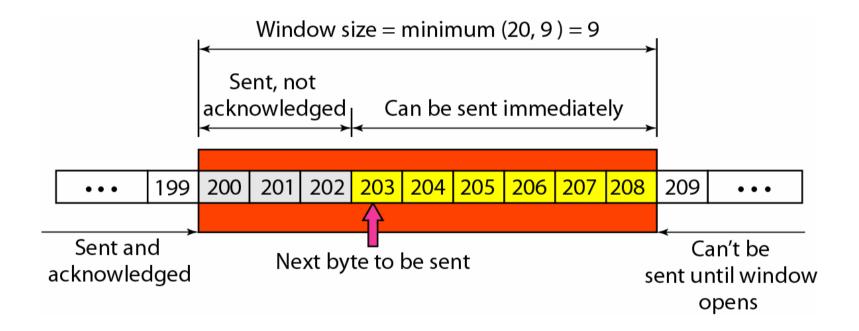
Solution

The size of the window is the smaller of rwnd and cwnd, which is 3000 bytes.



Figure 23.23 shows an unrealistic example of a sliding window. The sender has sent bytes up to 202. We assume that cwnd is 20 (in reality this value is thousands of bytes). The receiver has sent an acknowledgment number of 200 with an rwnd of 9 bytes (in reality this value is thousands of bytes). The size of the sender window is the minimum of rwnd and cwnd, or 9 bytes. Bytes 200 to 202 are sent, but not acknowledged. Bytes 203 to 208 can be sent without worrying about acknowledgment. Bytes 209 and above cannot be sent.

Figure 23.23 *Example 23.6*





Some points about TCP sliding windows:

- □ The size of the window is the lesser of rwnd and cwnd.
- □ The source does not have to send a full window's worth of data.
- □ The window can be opened or closed by the receiver, but should not be shrunk.
- □ The destination can send an acknowledgment at any time as long as it does not result in a shrinking window.
- □ The receiver can temporarily shut down the window; the sender, however, can always send a segment of 1 byte after the window is shut down.



ACK segments do not consume sequence numbers and are not acknowledged.



In modern implementations, a retransmission occurs if the retransmission timer expires or three duplicate ACK segments have arrived.



No retransmission timer is set for an ACK segment.



Data may arrive out of order and be temporarily stored by the receiving TCP, but TCP guarantees that no out-of-order segment is delivered to the process.

Figure 23.24 Normal operation

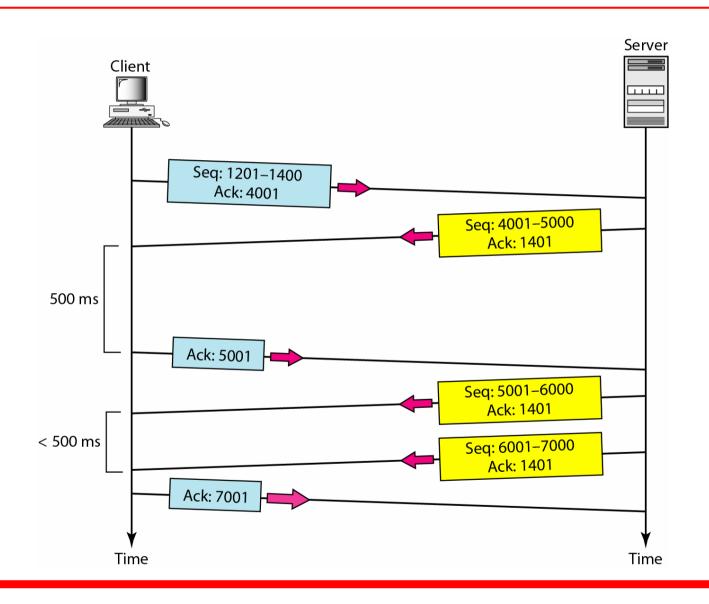
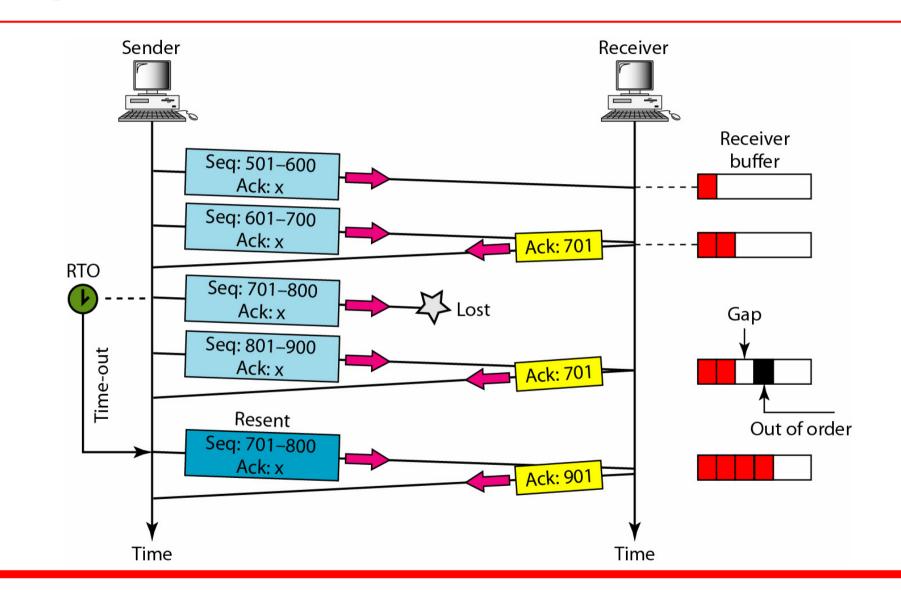


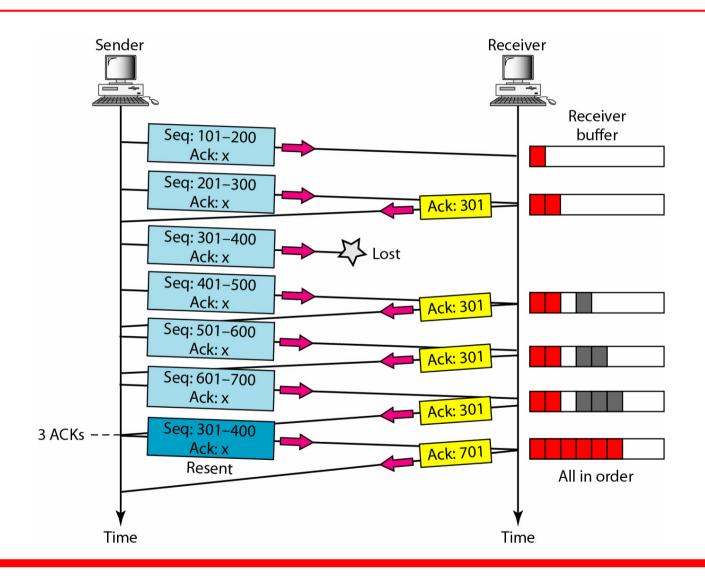
Figure 23.25 Lost segment





The receiver TCP delivers only ordered data to the process.

Figure 23.26 Fast retransmission



23-4 SCTP

Stream Control Transmission Protocol (SCTP) is a new reliable, message-oriented transport layer protocol. SCTP, however, is mostly designed for Internet applications that have recently been introduced. These new applications need a more sophisticated service than TCP can provide.

Topics discussed in this section:

SCTP Services and Features
Packet Format
An SCTP Association
Flow Control and Error Control

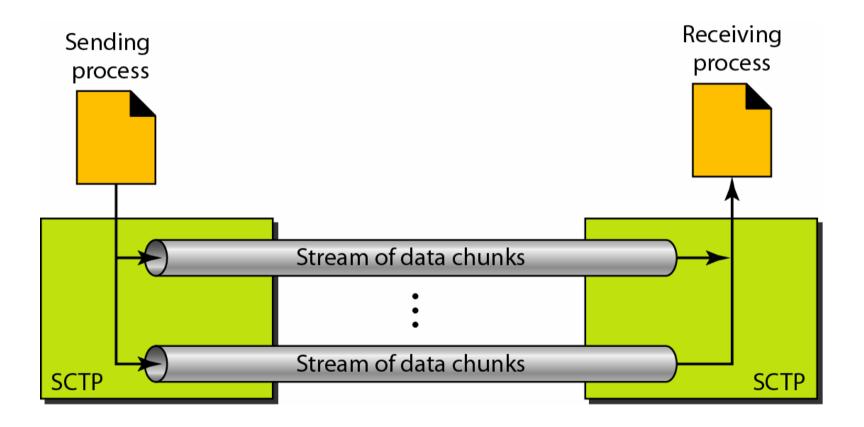


SCTP is a message-oriented, reliable protocol that combines the best features of UDP and TCP.

Table 23.4 Some SCTP applications

Protocol	Port Number	Description
IUA	9990	ISDN over IP
M2UA	2904	SS7 telephony signaling
M3UA	2905	SS7 telephony signaling
H.248	2945	Media gateway control
H.323	1718, 1719, 1720, 11720	IP telephony
SIP	5060	IP telephony

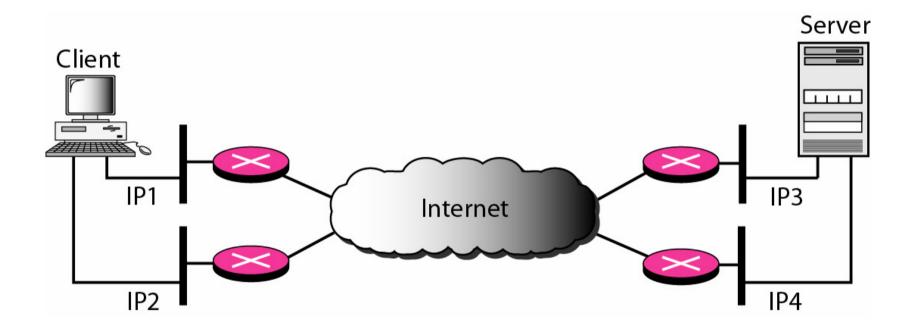
Figure 23.27 Multiple-stream concept





An association in SCTP can involve multiple streams.

Figure 23.28 Multihoming concept





SCTP association allows multiple IP addresses for each end.



In SCTP, a data chunk is numbered using a TSN.



To distinguish between different streams, SCTP uses an SI.

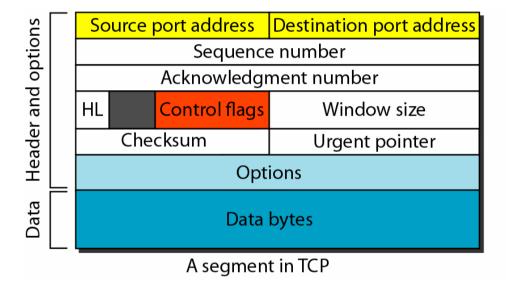


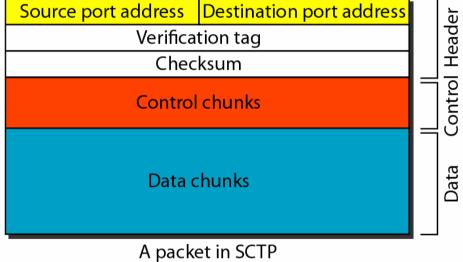
To distinguish between different data chunks belonging to the same stream, SCTP uses SSNs.



TCP has segments; SCTP has packets.

Figure 23.29 Comparison between a TCP segment and an SCTP packet

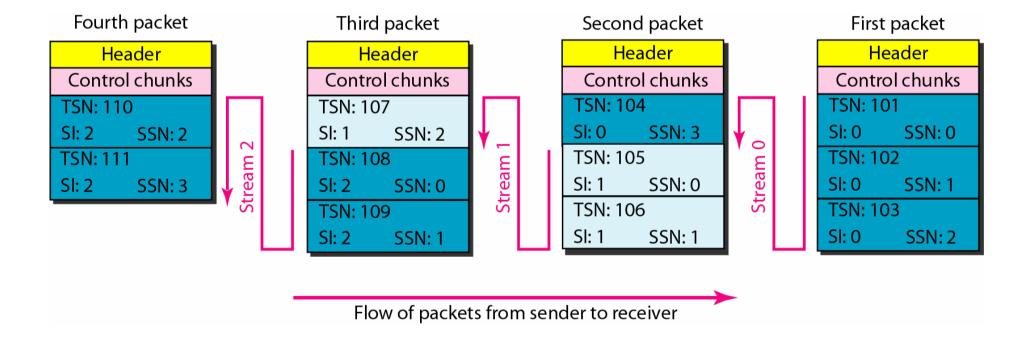






In SCTP, control information and data information are carried in separate chunks.

Figure 23.30 Packet, data chunks, and streams



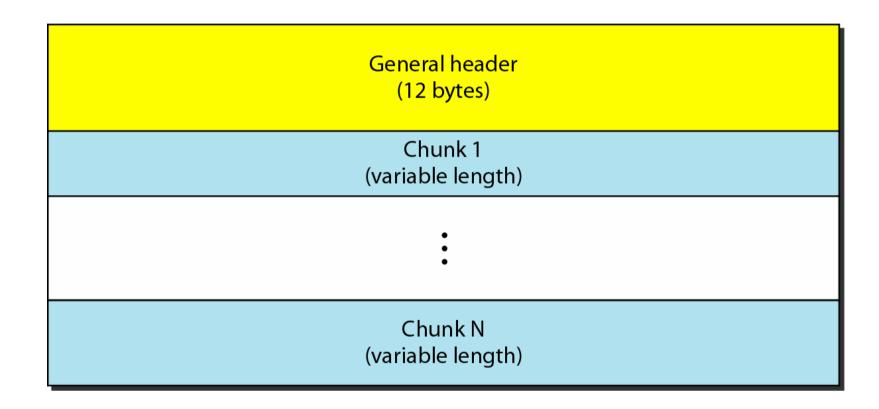


Data chunks are identified by three items: TSN, SI, and SSN.
TSN is a cumulative number identifying the association; SI defines the stream; SSN defines the chunk in a stream.



In SCTP, acknowledgment numbers are used to acknowledge only data chunks; control chunks are acknowledged by other control chunks if necessary.

Figure 23.31 SCTP packet format





In an SCTP packet, control chunks come before data chunks.

Figure 23.32 General header

Source port address 16 bits	Destination port address 16 bits	
Verification tag 32 bits		
Checksum 32 bits		

Table 23.5 Chunks

Туре	Chunk	Description
0	DATA	User data
1	INIT	Sets up an association
2	INIT ACK	Acknowledges INIT chunk
3	SACK	Selective acknowledgment
4	HEARTBEAT	Probes the peer for liveliness
5	HEARTBEAT ACK	Acknowledges HEARTBEAT chunk
6	ABORT	Aborts an association
7	SHUTDOWN	Terminates an association
8	SHUTDOWN ACK	Acknowledges SHUTDOWN chunk
9	ERROR	Reports errors without shutting down
10	COOKIE ECHO	Third packet in association establishment
11	COOKIE ACK	Acknowledges COOKIE ECHO chunk
14	SHUTDOWN COMPLETE	Third packet in association termination
192	FORWARD TSN	For adjusting cumulative TSN

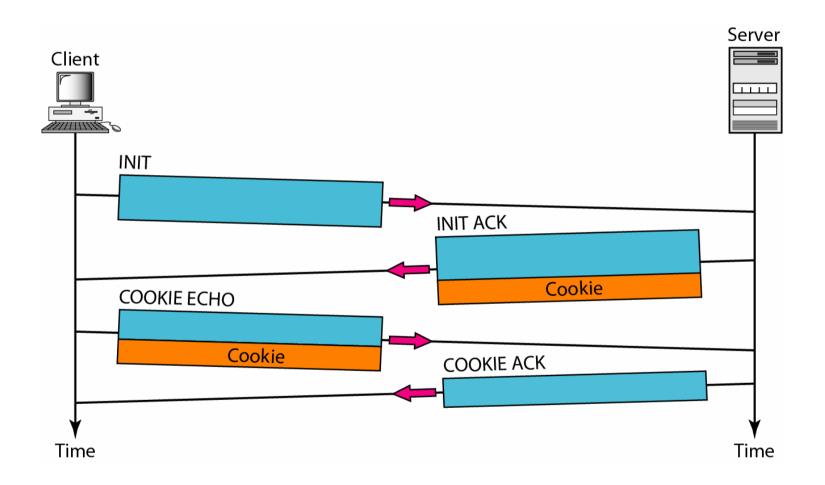


A connection in SCTP is called an association.



No other chunk is allowed in a packet carrying an INIT or INIT ACK chunk. A COOKIE ECHO or a COOKIE ACK chunk can carry data chunks.

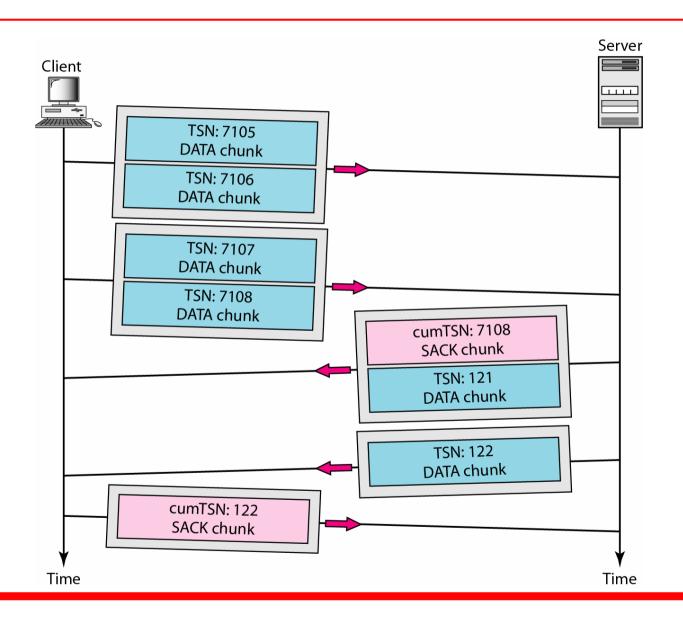
Figure 23.33 Four-way handshaking





In SCTP, only DATA chunks consume TSNs; DATA chunks are the only chunks that are acknowledged.

Figure 23.34 Simple data transfer





The acknowledgment in SCTP defines the cumulative TSN, the TSN of the last data chunk received in order.

Figure 23.35 Association termination

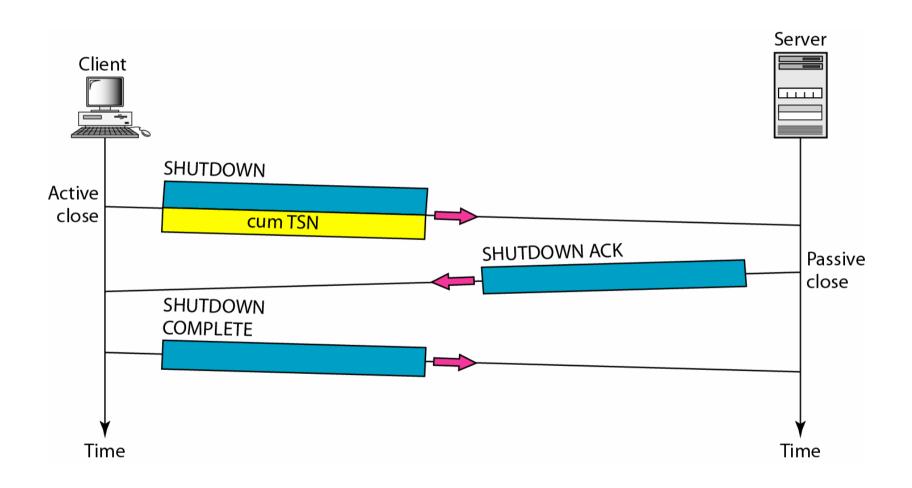


Figure 23.36 Flow control, receiver site

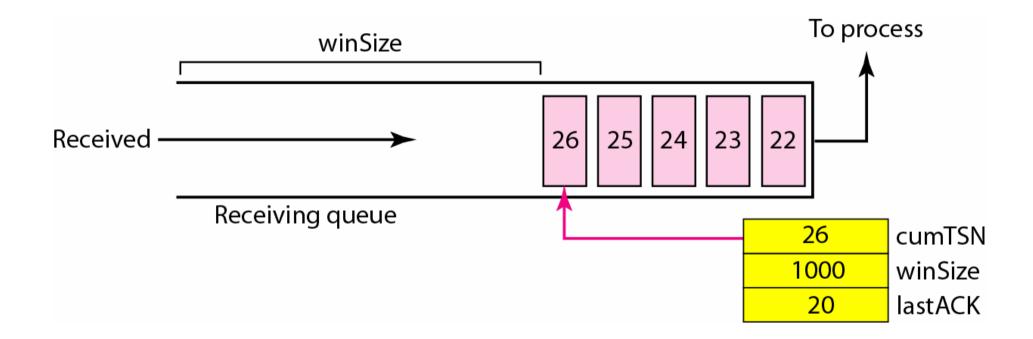


Figure 23.37 Flow control, sender site

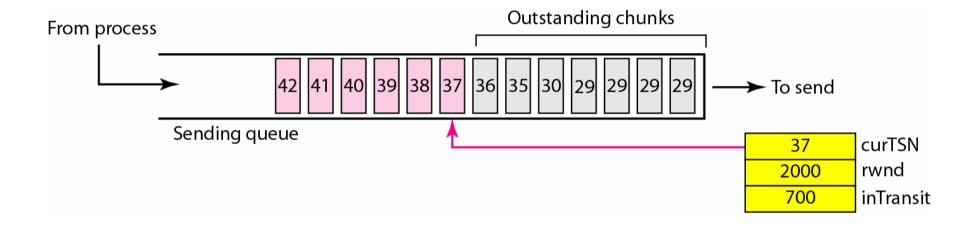


Figure 23.38 Flow control scenario

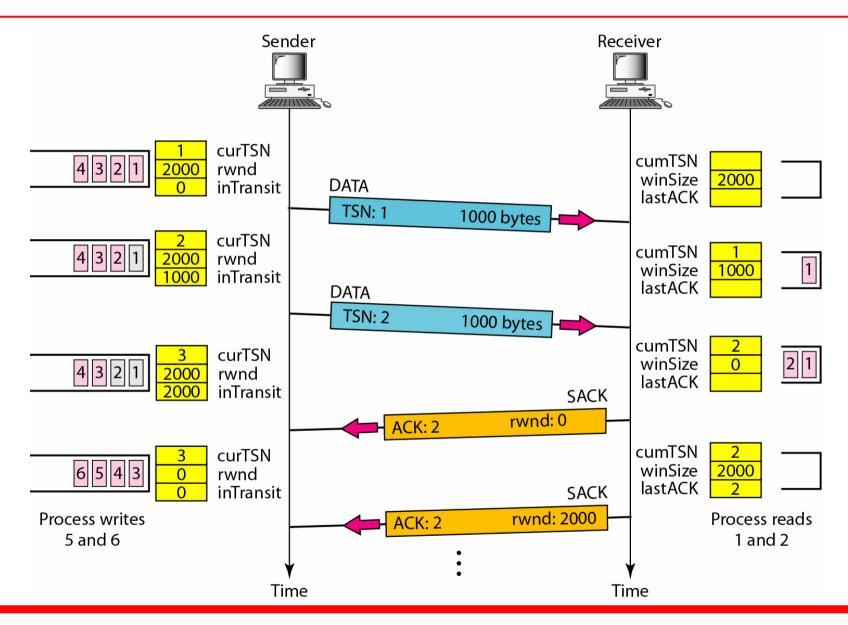


Figure 23.39 Error control, receiver site

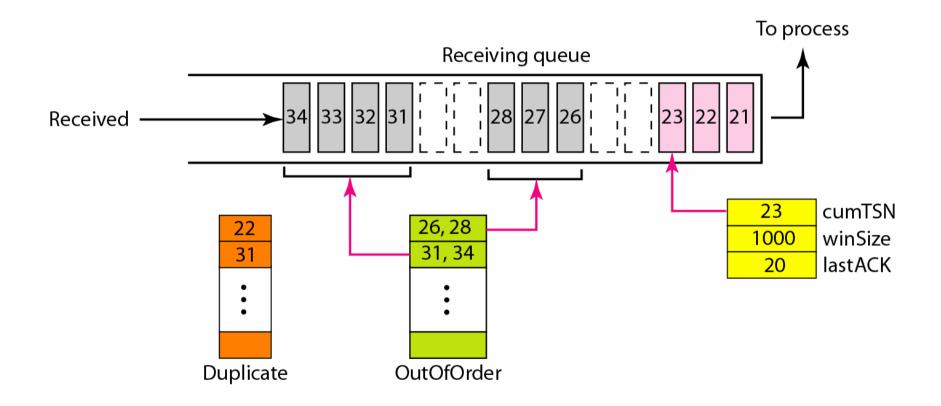


Figure 23.40 Error control, sender site

