## **Chapter 4**

# Process-to-Process Delivery: UDP and TCP

#### 4-1 PROCESS-TO-PROCESS DELIVERY

The transport layer is responsible for process-to-process delivery—the delivery of a packet, part of a message, from one process to another. Two processes communicate in a client/server relationship, as we will see later.

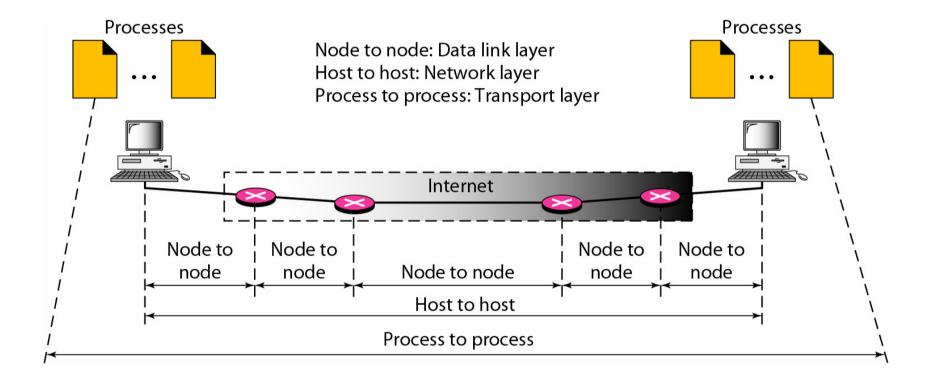
#### Topics discussed in this section:

Client/Server Paradigm
Multiplexing and Demultiplexing
Connectionless Versus Connection-Oriented Service
Reliable Versus Unreliable
Three Protocols

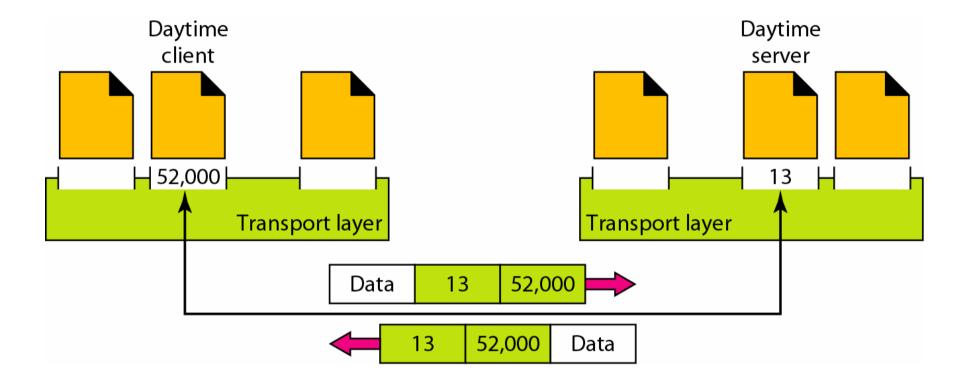


The transport layer is responsible for process-to-process delivery.

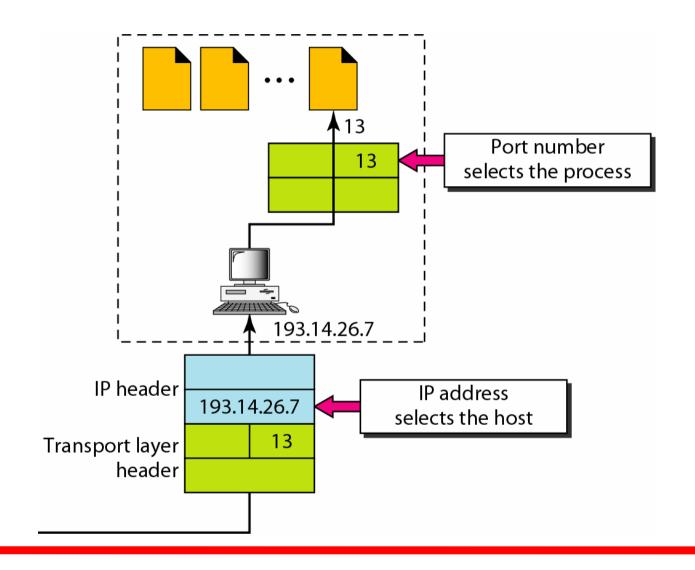
#### Figure 4.1 Types of data deliveries



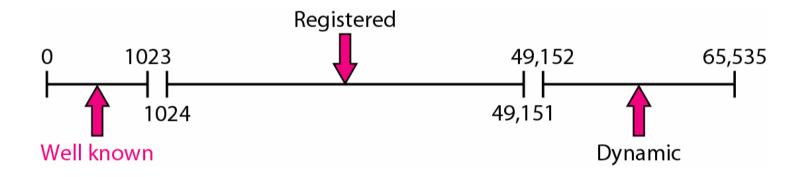
#### Figure 4.2 Port numbers



#### Figure 4.3 IP addresses versus port numbers



#### Figure 4.4 IANA ranges



#### Figure 4.5 Socket address

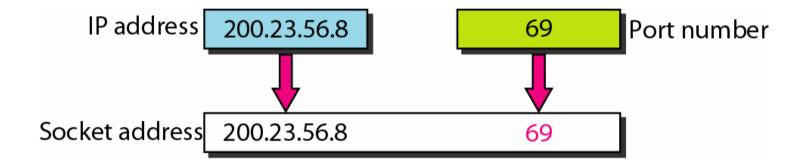
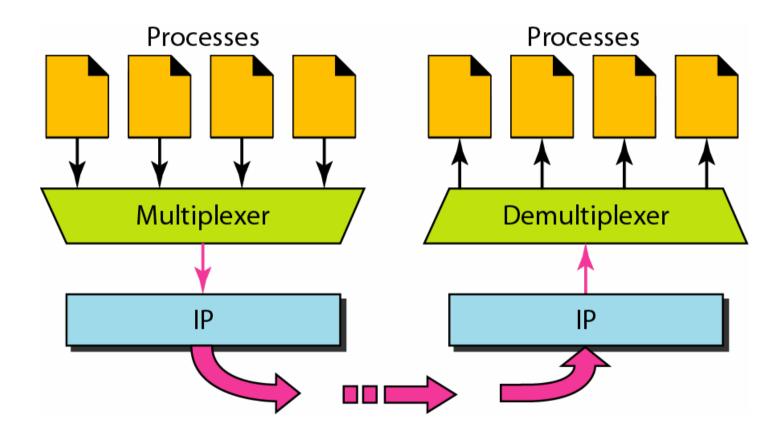


Figure 4.6 Multiplexing and demultiplexing



#### Figure 4.7 Error control

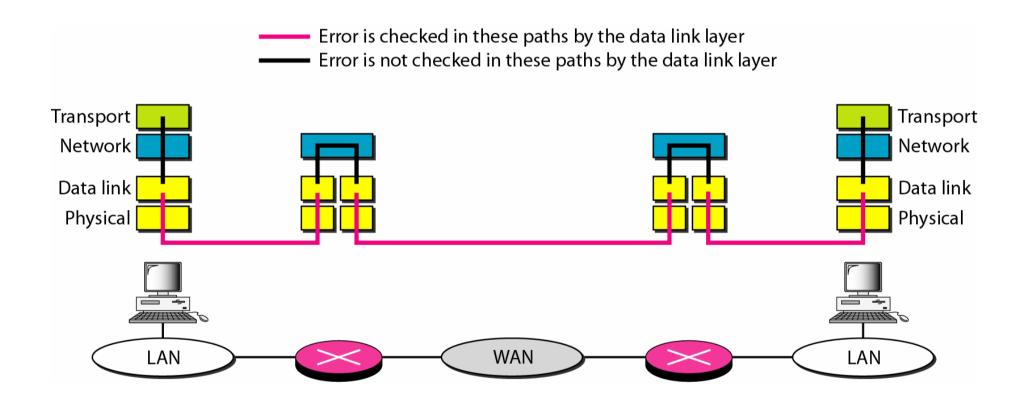
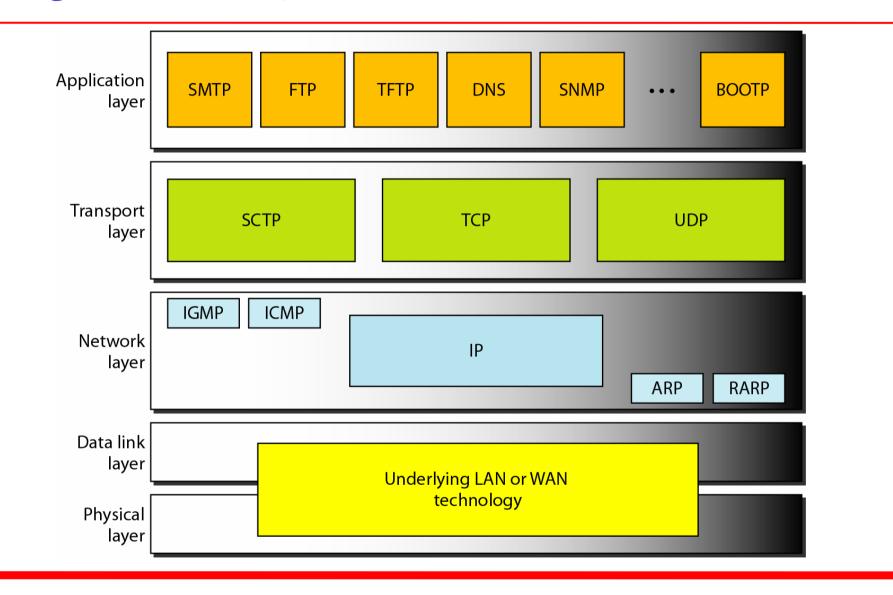


Figure 4.8 Position of UDP, TCP, and SCTP in TCP/IP suite



#### 4-2 USER DATAGRAM PROTOCOL (UDP)

The User Datagram Protocol (UDP) is called a connectionless, unreliable transport protocol. It does not add anything to the services of IP except to provide process-to-process communication instead of host-to-host communication.

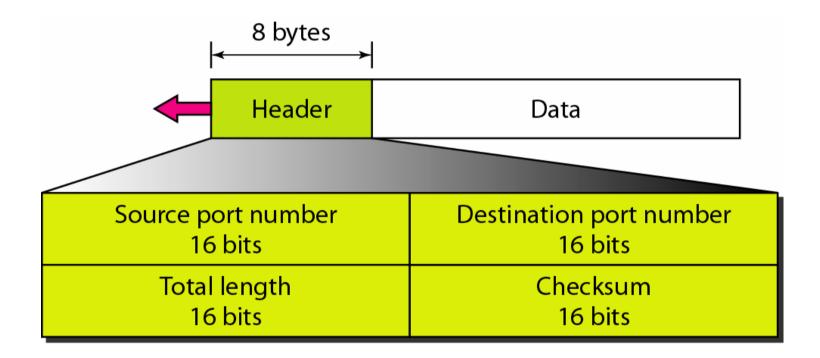
#### Topics discussed in this section:

Well-Known Ports for UDP
User Datagram
Checksum
UDP Operation
Use of UDP

 Table 4.1
 Well-known ports used with UDP

Port	Protocol	Description
7	Echo	Echoes a received datagram back to the sender
9	Discard	Discards any datagram that is received
11	Users	Active users
13	Daytime	Returns the date and the time
17	Quote	Returns a quote of the day
19	Chargen	Returns a string of characters
53	Nameserver	Domain Name Service
67	BOOTPs	Server port to download bootstrap information
68	BOOTPc	Client port to download bootstrap information
69	TFTP	Trivial File Transfer Protocol
111	RPC	Remote Procedure Call
123	NTP	Network Time Protocol
161	SNMP	Simple Network Management Protocol
162	SNMP	Simple Network Management Protocol (trap)

#### Figure 4.9 User datagram format

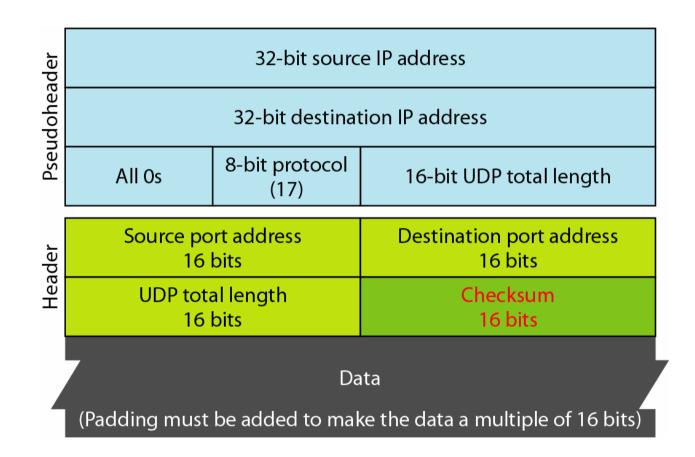




## **UDP** length

= IP length – IP header's length

#### Figure 4.10 Pseudoheader for checksum calculation



## Example 23.2

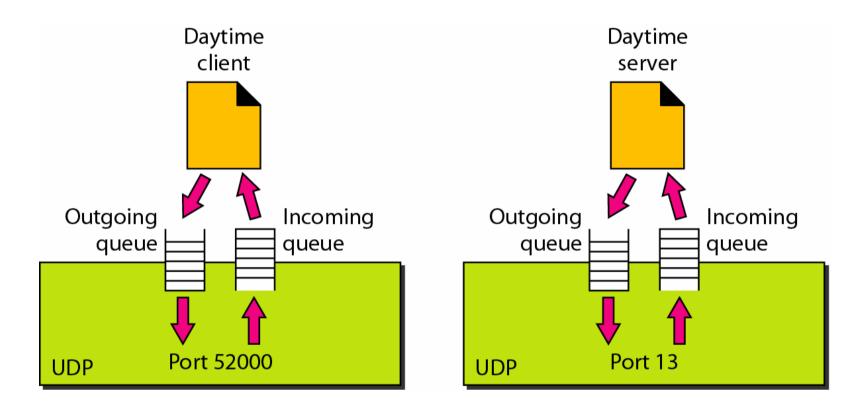
Figure 4.11 shows the checksum calculation for a very small user datagram with only 7 bytes of data. Because the number of bytes of data is odd, padding is added for checksum calculation. The pseudoheader as well as the padding will be dropped when the user datagram is delivered to IP.

#### Figure 4.11 Checksum calculation of a simple UDP user datagram

153.18.8.105						
171.2.14.10						
All Os 17		15				
10	87	13				
1	5	All Os				
Т	E	S	Т			
I	N	G	All Os			

10011001 00010010 — <b>&gt;</b> 153.18
00001000 01101001
10101011 00000010 — <b>&gt;</b> 171.2
00001110 00001010
00000000 00010001 $\longrightarrow$ 0 and 17
00000000 00001111 — <b>&gt;</b> 15
00000100 00111111
00000000 00001101 <del>&gt;</del> 13
00000000 00001111 <del></del>
0000000 00000000
01010100 01000101
01010011 01010100 <del>→</del> Sand T
01001001 01001110 → ➤ land N
01000111 00000000
<del></del>
10010110 11101011 <del>→</del> Sum
01101001 00010100 → Checksum

#### Figure 4.12 Queues in UDP





Connectionless Services: UDP provides a connectionless service. This means that each user datagram sent by UDP is an independent datagram



Flow and Error Control: UDP is a very unreliable transport protocol. There is no flow and error control.



Use of UDP: UDP is suitable for a process that requires simple request-response communication with little concern for flow and error control.

#### **4-3** TCP

TCP is a connection-oriented protocol; it creates a virtual connection between two TCPs to send data. In addition, TCP uses flow and error control mechanisms at the transport level.

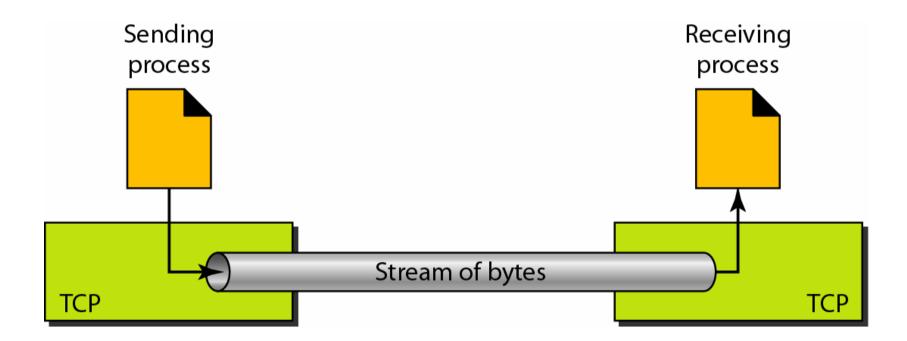
#### Topics discussed in this section:

TCP Services
TCP Features
Segment
A TCP Connection
Flow Control
Error Control

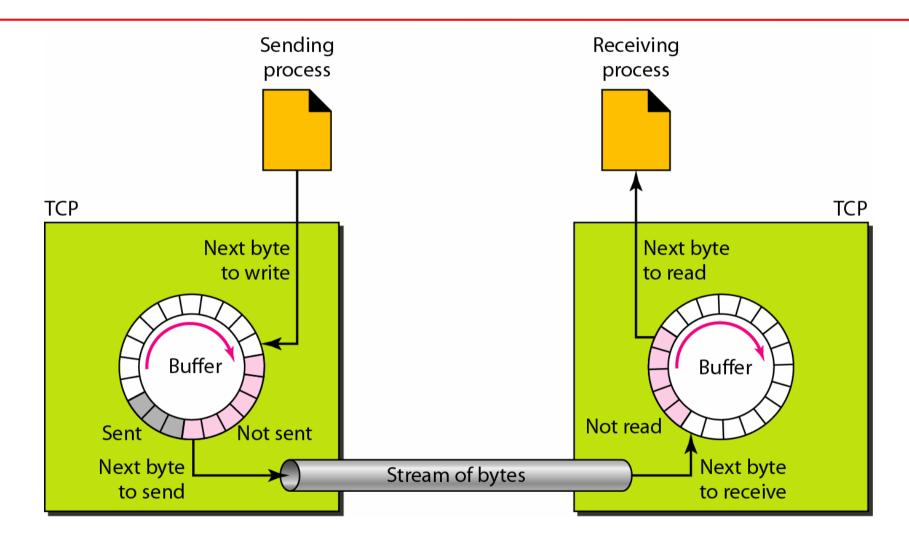
 Table 4.2
 Well-known ports used by TCP

Port	Protocol	Description
7	Echo	Echoes a received datagram back to the sender
9	Discard	Discards any datagram that is received
11	Users	Active users
13	Daytime	Returns the date and the time
17	Quote	Returns a quote of the day
19	Chargen	Returns a string of characters
20	FTP, Data	File Transfer Protocol (data connection)
21	FTP, Control	File Transfer Protocol (control connection)
23	TELNET	Terminal Network
25	SMTP	Simple Mail Transfer Protocol
53	DNS	Domain Name Server
67	ВООТР	Bootstrap Protocol
79	Finger	Finger
80	HTTP	Hypertext Transfer Protocol
111	RPC	Remote Procedure Call

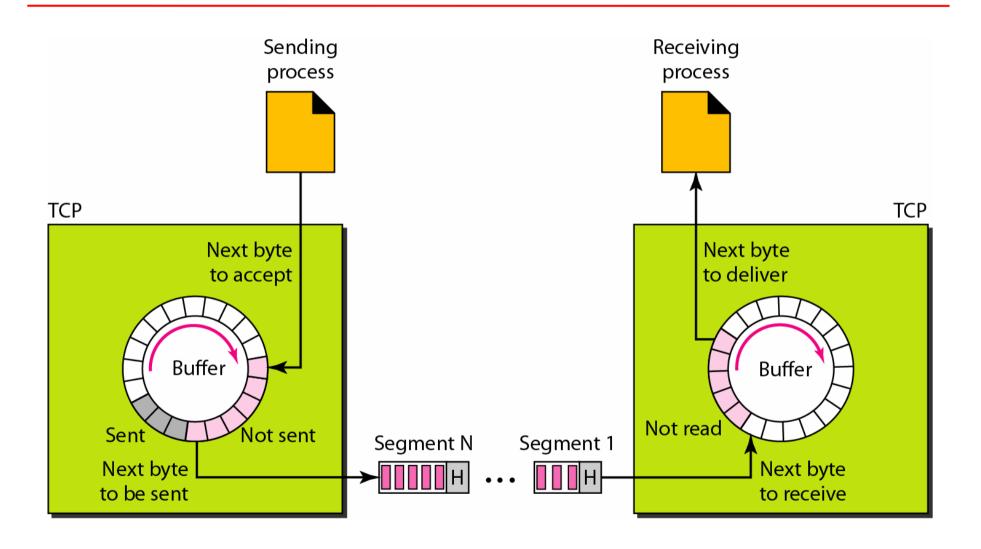
#### Figure 4.13 Stream delivery



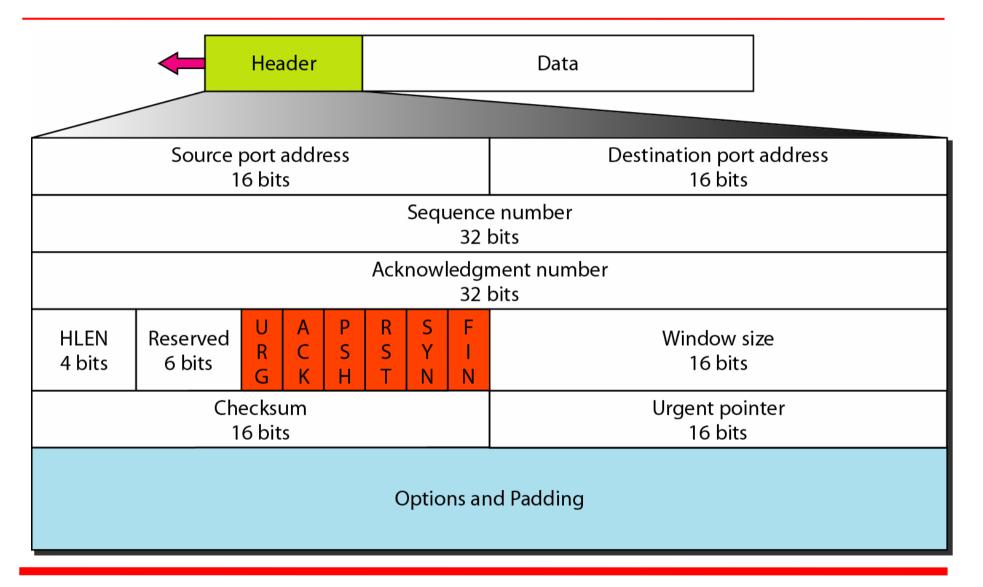
#### Figure 4.14 Sending and receiving buffers



#### Figure 4.15 TCP segments



#### Figure 4.16 TCP segment format





The bytes of data being transferred in each connection are numbered by TCP. The numbering starts with a randomly generated number.

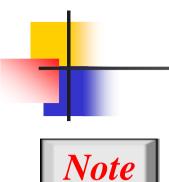


The value in the sequence number field of a segment defines the number of the first data byte contained in that segment.

## Example 4.3

The following shows the sequence number for each segment:

```
Segment 1 Sequence Number: 10,001 (range: 10,001 to 11,000)
Segment 2 Sequence Number: 11,001 (range: 11,001 to 12,000)
Segment 3 Sequence Number: 12,001 (range: 12,001 to 13,000)
Segment 4 Sequence Number: 13,001 (range: 13,001 to 14,000)
Segment 5 Sequence Number: 14,001 (range: 14,001 to 15,000)
```



The value of the acknowledgment field in a segment defines the number of the next byte a party expects to receive.

The acknowledgment number is cumulative.

#### Figure 4.17 Control field

URG: Urgent pointer is valid

ACK: Acknowledgment is valid

PSH: Request for push

RST: Reset the connection

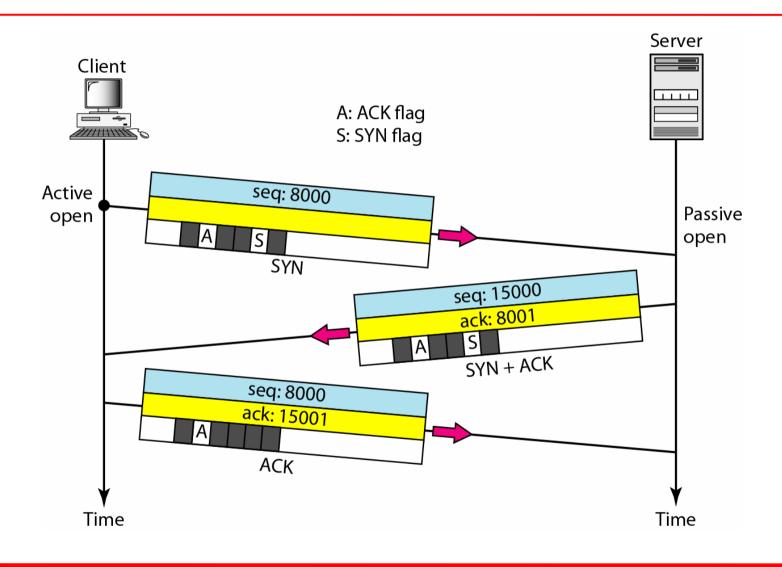
SYN: Synchronize sequence numbers

FIN: Terminate the connection

 Table 4.3 Description of flags in the control field

Flag	Description
URG	The value of the urgent pointer field is valid.
ACK	The value of the acknowledgment field is valid.
PSH	Push the data.
RST	Reset the connection.
SYN	Synchronize sequence numbers during connection.
FIN	Terminate the connection.

#### Figure 4.18 Connection establishment using three-way handshaking





## A SYN segment cannot carry data, but it consumes one sequence number.

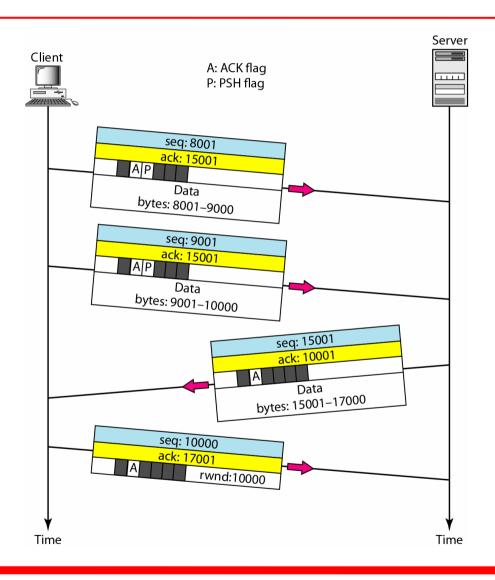


A SYN + ACK segment cannot carry data, but does consume one sequence number.

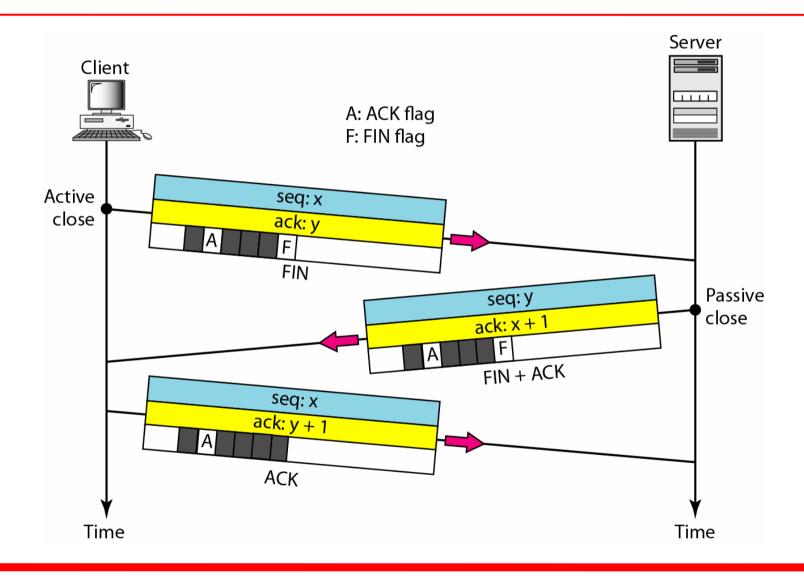


### An ACK segment, if carrying no data, consumes no sequence number.

#### Figure 4.19 Data transfer



#### Figure 4.20 Connection termination using three-way handshaking



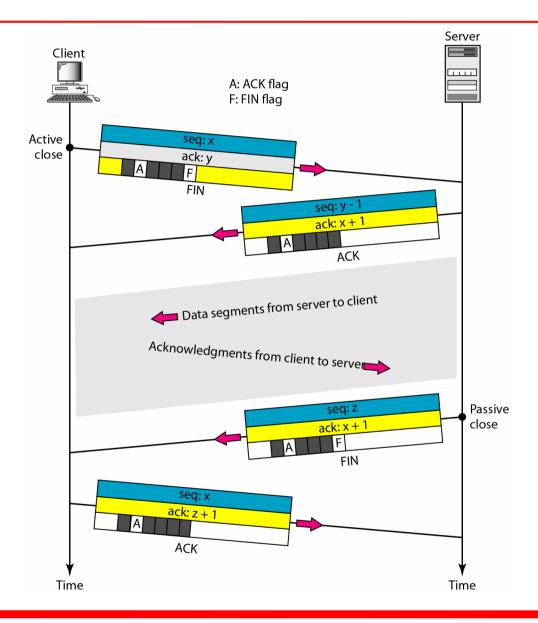


## The FIN segment consumes one sequence number if it does not carry data.



## The FIN + ACK segment consumes one sequence number if it does not carry data.

#### Figure 4.21 Half-close





Flow Control: TCP uses a sliding window to handle flow control. The sliding window protocol used by TCP, is something between the Go-Back-N and Selective Repeat sliding window.



A sliding window is used to make transmission more efficient as well as to control the flow of data so that the destination does not become overwhelmed with data.

TCP sliding windows are byte-oriented.

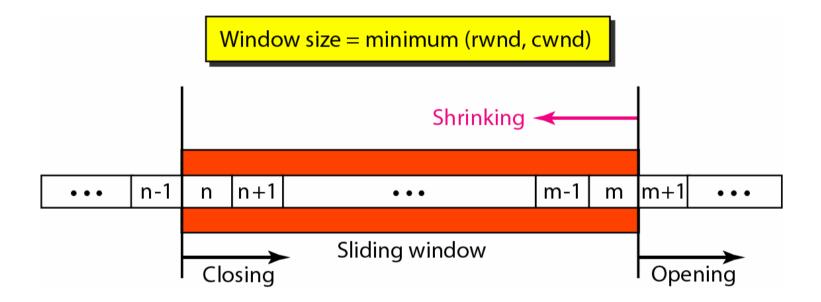


- TCP sliding windows are byte-oriented; the data link layer is frame-oriented.
- The TCP's sliding window is of variable size; the data link layer was of fixed size.



The size of the window at one end is determined by the lesser of two values: receiver window (rwnd) or congestion window (cwnd).

#### Figure 4.22 Sliding window



## Example 4.5

What is the size of the window for host A if the value of rwnd is 3000 bytes and the value of cwnd is 3500 bytes?

#### **Solution**

The size of the window is the smaller of rwnd and cwnd, which is 3000 bytes.



The receiver window is the value advertised by the opposite end in a segment containing acknowledgment. It is the number of bytes the other end can accept before its buffer overflows

## Example 4.4

What is the value of the receiver window (rwnd) for host A if the receiver, host B, has a buffer size of 5000 bytes and 1000 bytes of received and unprocessed data?

#### Solution

The value of rwnd = 5000 - 1000 = 4000. Host B can receive only 4000 bytes of data before overflowing its buffer. Host B advertises this value in its next segment to A.

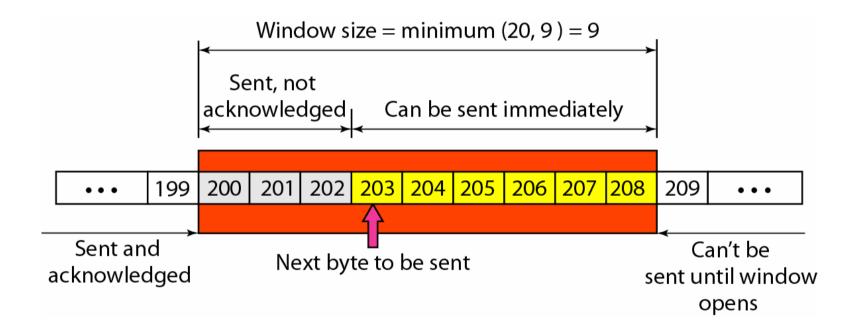


# The congestion window is a value determined by the network to avoid congestion.

## Example 4.6

Figure 4.23 shows an example of a sliding window. The sender has sent bytes up to 202. We assume that cwnd is 20. The receiver has sent an acknowledgment number of 200 with an rwnd of 9 bytes. The size of the sender window is the minimum of rwnd and cwnd, or 9 bytes. Bytes 200 to 202 are sent, but not acknowledged. Bytes 203 to 208 can be sent without worrying about acknowledgment. Bytes 209 and above cannot be sent.

#### Figure 4.23 Example 4.6





#### Some points about TCP sliding windows:

- The size of the window is the lesser of rwnd and cwnd.
- □ The source does not have to send a full window's worth of data.
- □ The window can be opened or closed by the receiver, but should not be shrunk.
- □ The destination can send an acknowledgment at any time as long as it does not result in a shrinking window.
- □ The receiver can temporarily shut down the window; the sender, however, can always send a segment of 1 byte after the window is shut down.



Error Control: Error detection and correction in TCP is achieved through the use of three simple tools: checksum, acknowledgment, and time-out.



#### **Checksum:**

Each segment includes a checksum field which is used to check for a corrupted segment.



# Acknowledgment (ACK) TCP uses acknowledgments to confirm the receipt of data segments. ACK segments do not consume sequence numbers and are not acknowledged.



## Retransmission: The heart of the error control mechanism is the retransmission of segments.

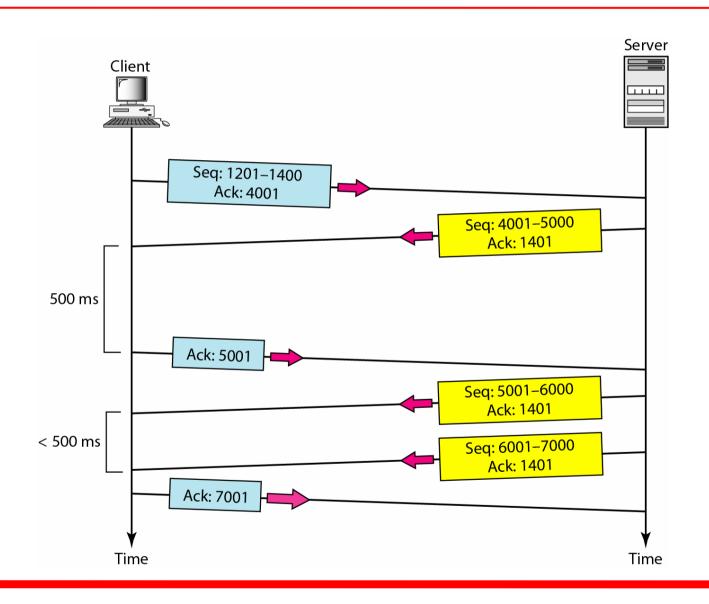


## Note: No retransmission timer is set for an ACK segment.

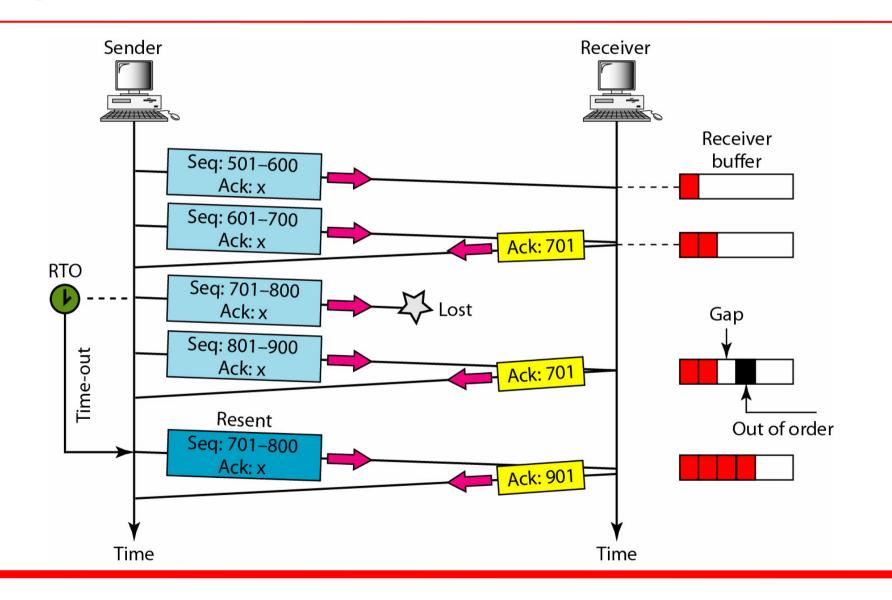


Data may arrive out of order and be temporarily stored by the receiving TCP, but TCP guarantees that no out-of-order segment is delivered to the process.

#### Figure 4.24 Normal operation



#### Figure 4.25 Lost segment





The receiver TCP delivers only ordered data to the process.

#### Figure 4.26 Fast retransmission

